

Kenyan News Monitoring Web Crawler for Social Harm Detection

Author & Developer: Mr. Patrick

Date: November 2025

1. Overview

- Scraping: Built a robust, multilingual (English + Swahili + Kikuyu-aware) news scraper targeting 15 major Kenyan news sites
- ML/NLP: Integrated Named Entity Recognition (NER) and multilingual sentiment analysis with keyword-based social-harm categorization (GBV, Cyberbullying, Scams)
- Backend: Designed and implemented a Supabase (PostgreSQL) backend with real-time dashboard auto-refresh via triggers
- Deployment: Created a fully automated nightly GitHub Actions pipeline running at 00:00 EAT every day

2. Objectives

The main goal of this project is to automatically monitor Kenyan online media every single day and detect articles related to three critical social harm areas:

1. Gender-Based Violence (GBV) & Femicide
2. Cyberbullying & Online Harassment
3. Scams & Financial Fraud

Why this matters:

- These issues are severely under-reported in structured data in Kenya
- Early detection from news can inform NGOs, researchers, law enforcement, and policymakers
- The system runs completely autonomously

3. Scraping / Data Collection

Libraries Used

selenium + webdriver-manager → for JavaScript-heavy sites

newspaper3k → best-in-class article extraction (title, text, date, authors)

beautifulsoup4 + lxml → link extraction from homepage

pandas → data structuring

nltk → sentence tokenization for keyword matching

transformers + torch → Hugging Face NER & sentiment models

supabase-py → direct database uploads

Scraping Workflow

1. Headless Chrome launches (stealth mode to avoid detection)
2. Visits each of the 15 target sites (Tuko, Citizen, Nation, The Star, etc. + Swahili & Kikuyu sections)
3. Waits for page to fully load (60-second timeout)
4. Extracts up to 10 latest article links per site using strict path filters (`/news/`, `/article/`, `/kenya/`, `/swahili/`, etc.)
5. For each link → `newspaper3k.Article()` downloads and parses clean text, title, and publish date
6. Filters:
 - Only articles < 30 days old
 - Minimum 150 characters of body text
7. Runs categorization + NLP analysis
8. Stores everything in a Pandas DataFrame → CSV + Supabase

Challenges & Solutions

Challenge	Solution Implemented
Heavy JavaScript sites	Selenium with headless Chrome + long waits
Bot detection	Disabled automation flags, excluded switches, spoofed navigator.webdriver
Duplicate articles	`article_url` is UNIQUE in DB + upsert on conflict
robots.txt compliance	Respectful delays (2–10 sec), only public news sections, no aggressive crawling
Dynamic "infinite scroll"	Limited to top 10 visible links (most sites show latest on homepage)

4. Data Preprocessing / ML Pipeline

Categorization (Rule-Based + Multilingual Keywords)

I chose a hybrid keyword approach instead of full classifier training because:

- Very limited labeled training data in Swahili/Kikuyu
- Need for immediate interpretability and low latency

Machine Learning Components

Task	Model Used	Reason Chosen
Named Entity Recognition `	`Davlan/bert-base-multilingual-cased-ner-hrl`	Excellent performance on African names/locations
Sentiment Analysis	`cardiffnlp/twitter-xlm-roberta-base-sentiment`	Trained on multilingual tweets, understands code-switching

Both models run locally (cached via GitHub Actions) – no external API costs.

5. Backend & Storage

Database: Supabase (PostgreSQL as a service) – chosen for:

- Free tier sufficient
- Built-in REST API & Auth
- Real-time subscriptions
- Easy SQL editor & triggers

Main Tables

1. `scraped_articles` – raw article data (one row per article)
2. `dashboard` – one row per day with aggregated statistics (auto-refreshed)

Smart Dashboard System

I built a self-updating dashboard using PostgreSQL triggers:

- Every time an article is inserted/updated/deleted → trigger fires
- `refresh_dashboard_for_date(CURRENT_DATE)` recalculates ALL metrics for today
- Uses `ON CONFLICT (date) DO UPDATE` → always exactly one row per day

This means your dashboard is always 100% up-to-date without running separate jobs.

6. Deployment / Automation

Fully automated nightly run via GitHub Actions (`scraper.yml`)

- Schedule: Every day at 00:00 East Africa Time (21:00 UTC)
- Manual trigger available (`workflow_dispatch`)
- Caches: pip, Chrome, Hugging Face models → cold start < 2 min, warm < 40 sec
- Secrets: `SUPABASE_URL` and `SUPABASE_KEY` stored securely
- Artifact: Daily CSV downloadable from GitHub Actions UI

7. References / Links

- GitHub Repository: https://github.com/Muchai10/safeguard_crawler
- Main Script: `WebCrawler_V2.py`
- Database Schema & Dashboard Logic: `Supabase.sql`
- CI/CD Pipeline: `./github/workflows/scraper.yml`
- Requirements: `requirements.txt`
- Daily CSV Artifacts: Available in GitHub Actions → "scraped-articles.csv"